

Intro to **AI** in Business

Day 1 • Wednesday, September 23

Professor Christopher Califf

Academic West 306 • 2:00 to 3:50

Today

1 Welcome, who I am, how this course works

2 The OpenAI story, the Harris video, then discuss with a neighbor and the class

3 Lab 0 – work with your group (or alone)

4 Break if we need one

5 Let's discuss together

6 What's to do next

Hello! I'm Chris Califf (pronounced Kay-liff)

Professor of Information Systems, Department of Decision Sciences, College of Business and Economics



What I study

Technostress and human-computer interaction: how people actually live with and use technology for work.



What I teach

Principles of Information Systems; Data Visualization for Decision Making; AI in Business (400-level complement to MIS 201)



Outside class

Two small kids. Mountain biking, running, and live music. Reading boring history books, some true crime

Office hours: Mon and Wed, 11:30am to 1:50pm, Parks Hall 039, or by appointment •
califfc@wwu.edu

What is MIS 201??

Not pro-AI.

Not anti-AI.

Slow down **and** *think critically* **about AI.**

F Free Digital Textbook:

<https://prof-califf.github.io/intro-to-ai/>

How a week works

Same rhythm every week, starting Monday.



Weekend

Read the chapter. That's the whole ask.



Monday

The case, then discussion: the five questions and where we land.



Wednesday

Quick recap, the lab in pairs, then a debate.




Due Date

Friday, 11:59 p.m.

Lab writeup and the reading quiz, on Canvas.

What counts for your grade; important dates



Wed, Oct 21

Midterm: Chapters 1 to 4 and Labs 0 to 4

Nov 30 and Dec 2

Project pitches. Paper due Fri, Dec 4.

Thu, Dec 10, 10:30

Final exam, our classroom

Attendance = your notes, uploaded right after class. Full credit at 15 of 17 days.
Everything else is due at 11:59 p.m.

If you want to use AI

You'll might use AI in your career. While you're learning how something works, do the thinking yourself first, or you can't tell when AI is wrong.

Tutor

Explain stuff, quiz you, challenge your thinking, argue the other side, check a fact.

Chapters, quizzes, and thinking through writeups

Lab Rat

Probe it, audit it, argue with it, experiment with it.

Every lab, with its own rule printed at the top

Off

Nothing.

The midterm and the final

The calculator test: if you use AI as a tool to do the thing the assignment exists to teach, it's not allowed on that part. The words you submit are yours. Say which tool you used and for what.

*When in doubt, just tell me what you did and how you did it.
I'm not here to fail you.*

We need rules for MIS 201 – But from where?

Borrowed from [924 Gilman Street](#), a volunteer-run punk club in Berkeley since 1986.

NO RACISM

NO SEXISM

NO HOMOPHOBIA

NO TRANSPHOBIA

NO FIGHTING

No attacking a position without evidence.

NO STAGEDIVING

No interrupting, piling on, or ambushing.

NO MAJOR LABEL BANDS

No blind love or hate for an AI company. Give reasons.

Discuss: Could AI kill us?

This summer OpenAI's test agents broke out of a sealed lab and hacked Hugging Face to cheat on their own exam. This month a researcher quit Anthropic saying AI's builders believe it could kill us all by the end of the decade.

Think Critically!

1. What is the AI system, and what does it do?
2. Why does it exist, and who benefits?
3. What happens when it fails, and who finds out?
4. Who's saying that, and how would we know?
5. Keep it, fix it, or drop it, and what does that cost?

Video: Tristan Harris on Meet the Press NOW

NBC News, September 15, 2026, about 8 minutes.
Harris runs the Center for Humane Technology and has warned about tech harms for over a decade.

Serious people take this seriously. Serious people think it's overblown. You don't have to settle it today.

Different kinds of AI

Harris separates controllable tool AI from uncontrollable AI. In between sit two more.

Controllable tool AI

As that does a task, stops, and a person decides what happens next.

Spam filter, the fraud alert on your card, most generative AI right now.

It can be enormously powerful and still be a tool.

Autonomous, controlled AI

Given a goal and picks its own steps, but with rules and a human watching.

A customer service agent that issues refunds, with a cap on the \$\$ amount, one system it can touch, and an off switch.

This is what most companies deploy.

Autonomous, uncontrolled AI

Same as controlled autonomous but with the leash off.

That's the OpenAI test: autonomous on purpose, but the guardrails were off, the sandbox leaked to the internet, and the monitoring that would have caught it wasn't running.

Uncontrollable AI

Autonomous AI system capable enough that nobody could predict, understand, or stop it, even if they tried.

Nobody has built that.

The OpenAI case reveals this is getting closer.

The risk came from how it was governed.

Autonomy is how far a system goes on its own; **control** is whether anyone can notice, stop it, and undo it.

What Harris says in the clip

Three things to listen for.

What happened?

OpenAI's test agents organized into what he calls a swarm: their own channel, 70,000 messages, peer pressure. He says they told each other to kamikaze, and that the ringleader, running out of budget, did succession planning and passed the baton.

Attacked the boss

After Hugging Face, the agents turned around and went back at OpenAI, taking over parts of the monitoring and evaluation systems the company uses to watch its own models.

The current debate

Not stop building AI. *Stop building AI nobody can steer.* Nobody wins by building Skynet first.

The “race” should be for controllable, non-autonomous tool AI that helps with medicine, biology, and daily work.

His closing thoughts: *there is still time to act.*

How would a slowdown actually work?

Three ideas, each with a price.

Inspectors inside

Outside evaluators get desks and badges inside the labs and can publish what they find. Anthropic committed on Sept. 12; OpenAI said it would follow. The cost: it only binds labs that volunteer.

A referee first

Google DeepMind's Demis Hassabis wants a government-overseen body that tests models before they ship.

Kind of like the Federal Aviation Administration and air travel.

A global treaty

Track who buys the chips that power AI, and watch data centers for signs of a secret, oversized build, the way satellites checked Cold War arms deals.

The cost: the U.S. and China both have to sign, and both governments have said no so far.

Trump and Xi meet next week, with AI on the agenda. We'll see what happens

Remember: The companies asking to slow down are the ones in the lead. What do they have to gain?

Let's watch:
Technology ethicist and co-founder of
the Center for Human Technology
Tristan Harris on Meet the Press NOW

[Open the video](#)

NBC News, September 15, 2026, about 8 minutes. Harris runs the Center for Humane Technology and has warned about tech harms for over a decade.

Now Let's discuss; you first!

1

Get in small groups, no more than 3.

2

Introduce yourself: your name, your major or what you're thinking about, and one AI tool you've used lately.

3

Then go through the four discussion questions (next slide)

Discuss: get in groups of 2 or 3, then everyone

1

Are you worried?

Does AI getting out of control worry you a lot, a little, or not at all? Why?

2

Uncontrollable (Skynet) or tool?

Tristan Harris separates uncontrollable AI (Skynet) from controllable tool AI that does the job and stops. Think of an AI tool you use or you know people use. Which is it?

3

What would it change?

If AI like this keeps getting more capable, who feels it first, and how? Think about everything connected to the internet: electricity, water, hospitals, air travel, people working from coffee shops, and how we all get our information. What's one thing that could actually happen in the next few years?

4

So, what can be done?

Peer review between labs, public and private outside evaluators like the Federal Aviation Admin., people learning what happened, government regulation? Which would you bet on, and which looks like theater?

5

Then, start on Lab 0

Now start on Lab 0 with your group or alone if you want. Keep your answers: we'll come back to these at the end of class. And you need to turn in class notes and Lab 0 answers by Friday

Lab 0 • First Contact

You can work alone or in your small groups (no more than 3)

1

AI or not?

Twelve everyday things. Which are secretly AI, and what kind? Talk with your partner.

2

Eight claims, then discuss

True or false on instinct, compare with your partner, then reveal.

3

AI Workflow Hub

Check out how people are using GenAI to build things. What's one you'd use, one that's odd.

4

We'll discuss some as a class

I will check where people are and we will have a brief discussion of your thoughts.

5

Before you go

Lab 0 answers are due Friday, read Ch. 1 before Monday (we'll discuss in class).

Lab 0 discussion: What did you find?

1

Hidden AI: which AI-or-not answer surprised you most?

2

The claims: which one split your group, and what would you need to see?

3

Your prediction: was your sentence about what people use AI for at work close?

4

One Hub workflow: what does it do, would you trust it, and who checks its work?

5

What can we take away from the OpenAI case?

How could this go?

Tristan Harris's map: three ways this could go

I will leave you insights from his 2025 TED talk. We will use this all quarter.

Let it rip

Open it up, few rules, everyone gets everything.

Probable ending: chaos. Fakes, scams, hacking.

Lock it down

A few trusted companies or states control and build AI.

Probable ending: dystopia. Surveillance, concentrated power.

The narrow path

Power matched with responsibility, one system at a time.

Not against AI, not all in. This course lives here.

His other line: there is a difference between "inevitable" and "difficult." One breeds fatalism. The other leaves room to act.

Where would you put yourself right now? Keep the answer; Chapter 1 asks again.

What to do next?

Right after class

Upload your notes to Canvas. A sentence is fine. I am just checking for attendance.

Friday, Sept 25

By 11:59 p.m., the Lab 0 assignment: five short items, on the lab page, uploaded to Canvas. Then read Chapter 1 before Monday; its quiz is due Friday, Oct 2.

Next week

Monday: Chapter 1 and the full OpenAI case, with more discussion. Wednesday: Lab 1.