
MIS 201 • WEEK 2 • WEDNESDAY 9/30

Lab 1: Inside the Machine

Lab first, then the debate. Groups of 2 or 3.

Today

2:00	Setup/Recap from Monday	5 min
2:05	Lab 1, sections 1 to 4	40 min
2:45	Recap: what did you just see?	10 min
2:55	Break	10 min
3:05	Debate (follow the Gilman Rules!)	35 min
3:40	Wrap-up and attendance	10 min

DUE FRIDAY 10/2, 11:59 P.M.

- Lab 1 writeup, five items
- Chapter 1 quiz

Answer the writeup items as you go today, and you'll be mostly done by 3:50.

Recap from Monday

1

Three kinds of AI

Predictive judges, generative creates, embodied and agentic acts.

2

Does it stop?

Autonomy is how far it goes on its own. Control is whether anyone can notice, stop, and undo it.

3

The five parts

Data, model, output, feedback loop, and humans: the part the marketing leaves off.

4

The case

Built as tool AI, ran for two months as uncontrolled AI. Hugging Face found out first; OpenAI learned from them.

5

Where we landed

Keep the testing, fix the containment. The fight is over who checks, and who pays for it.

WHAT DO YOU THINK?

Explain the OpenAI case to someone next to you in one sentence. Use at least one idea from this slide. Go.

Lab 1: Today you try two tiny AIs and watch an agent, then argue who should have been told.

Part 1: the lab

40 MIN · GROUPS OF 2 OR 3

1 Train an AI to see 10 min **done by 2:20**

2 Train an AI to play 10 min **done by 2:30**

3 Watch an AI agent 10 min **done by 2:40**

4 Under the hood 10 min **done by 2:50**

HOW IT WORKS

Work alone or in a group of 2 or 3.
One machine is fine.

Follow the steps on the page, in
order.

Take the screenshots the page asks
for.

**Answer the writeup items as you
go.**

Recap: what did you just see?

1

Mini-ImageNet

Did the ImageNet AI get anything wrong, right, and why?

2

Mini-AlphaGo

Nobody told it a strategy. Did it get better? How?

3

The AI agent, no human

When the robot got out of the sandbox, was it trying to escape? What was it actually doing?

4


The AI agent, regulation? Safety?

In your runs, what kept the robot in: a human saying no, a human on the kill switch, or a patched wall? If you were writing a rule for real AI labs, which one would you require, and who would check that they did it?

5

AI agents and you

Would you trust an AI agent to act for you, like booking travel or spending money? What would it need to have first?



BREAK • 5 MINUTES

Back in five

When you come back: the Gilman Street rules, then get into a group of 2 or 3 and pick a side.

Debate rules: the Gilman Street rules

APPLY THESE TODAY

Same rules as Week 1, borrowed from 924 Gilman Street. Here's what each one means in a debate.

**NO RACISM · NO SEXISM · NO
HOMOPHOBIA · NO TRANSPHOBIA**

Argue with ideas, never at people.

NO FIGHTING

No attacking a position without evidence. "That's wrong" isn't a point. "That's wrong, because Hugging Face found it first" is.

NO STAGEDIVING

No interrupting, piling on, or ambushing. One speaker at a time; let them finish. Respond to what they said, not to the person.

NO MAJOR LABEL BANDS

No blind love or hate for an AI company. OpenAI, Anthropic, Hugging Face: give reasons, not fandom.

The debate

35 MIN

When a safety test gets out of hand, who has to be told?

SIDE A

Disclosure should be required

Labs report breakouts and findings within days, to a regulator that can publish them.

Cost to name: every disclosure can teach attackers, and a reporting clock can slow fixes.

SIDE B

Voluntary, with an independent checker

Fund an outside auditor with real access, and let disclosure follow from it.

Cost to name: it took about a week, that was the best case, and nobody outside could check it until then.

The questions the debate turns on: the OpenAI and Hugging Face case

- 1 Who got hurt, and who paid to clean it up?
- 2 Who found out first, and how? Who should have?
- 3 If a rule required disclosure within days, what changes in July, and what does the rule cost?
- 4 Who should a report go to: the public, a regulator, an outside checker? What can each do with it?
- 5 What does your side's plan cost, and who pays?

Whatever side you pick, your points have to answer these.

How the debate runs: prep

16 MIN

Pick a side

2 min

1

Get into a group of 2 or 3. Pick Side A or Side B. If one side gets much bigger, I'll move a few groups over.

Research, with AI

10 min

2

Ask a chatbot for your side's best points, the other side's best objections, and evidence. Then close it and write your notes in your own words.

Team up and pick speakers

4 min

3

Join one or two other groups on your side, so your team is 6 or 7 people: three teams per side. Pick your single best point, an opener, and a rebutter.

YOUR NOTES

1. Your strongest point, in one sentence
2. The evidence behind it
3. The honest cost of your position
4. What the other side will probably say

THE RULE ON AI

Use it to research. Don't read from it.

How the debate runs: the rounds

17 MIN

Round 1

Openings

5 min

Openers to the front: A1, A2, A3, then B1, B2, B3. About 45 seconds each: your claim in one sentence, your evidence, and the cost of your plan.

Huddle

1 min

Back to your group. Rebutters pick one point the other side made.

Round 2

Rebuttals

5 min

Rebutters to the front. Side B first this time, then A. Answer one specific point, with evidence. No new arguments.

Round 3

Switch

3 min

30 seconds with your team: what was the other side's best point? Then someone new argues it in one sentence.

Recap

What did you hear?

3 min

Which argument from the other side was hardest to answer? Then we're done.

THE GILMAN RULES, APPLIED

One speaker at a time. Let them finish.

Every point needs evidence.

Answer the argument, not the person.

Reasons, not fandom.

No winner is declared. The point is to hear the best case on both sides.

BEFORE YOU LEAVE

Upload your notes

Attendance = your notes from today, uploaded to Canvas right after class.

At home: lab sections 6 and 7 (discussion, and the five questions on your own AI).

Friday 10/2, 11:59 p.m.: Lab 1 writeup and the Chapter 1 quiz, both on Canvas.