
MIS 201 • WEEK 2 • WEDNESDAY 9/30

Lab 1: Inside the Machine

Lab first, then the debate. Groups of 2 or 3.

Today

2:00	Setup/Recap from Monday	5 min
2:05	Lab 1, sections 1 to 4	40 min
2:45	Recap: what did you just see?	10 min
2:55	Break	10 min
3:05	Debate (follow the Gilman Rules!)	35 min
3:40	Wrap-up and attendance	10 min

DUE FRIDAY 10/2, 11:59 P.M.

- Lab 1 writeup, five items
- Chapter 1 quiz

Answer the writeup items as you go today, and you'll be mostly done by 3:50.

Recap from Monday

1

Three kinds of AI

Predictive judges, generative creates, embodied and agentic acts.

2

Does it stop?

Autonomy is how far it goes on its own. Control is whether anyone can notice, stop, and undo it.

3

The five parts

Data, model, output, feedback loop, and humans: the part the marketing leaves off.

4

The case

Built as tool AI, ran for two months as uncontrolled AI. Hugging Face found out first; OpenAI learned from them.

5

Where we landed

It's complicated. No clean keep, fix, or drop, and that's where today picks up.

WHAT DO YOU THINK?

Explain the OpenAI case to someone next to you in one sentence. Use at least one idea from this slide. Go.

Lab 1: Today you try two tiny AIs and watch an agent, then debate yesterday's AI accord.

Part 1: the lab

40 MIN · GROUPS OF 2 OR 3

1 Train an AI to see 10 min **done by 2:20**

2 Train an AI to play 10 min **done by 2:30**

3 Watch an AI agent 10 min **done by 2:40**

4 Under the hood 10 min **done by 2:50**

HOW IT WORKS

Work alone or in a group of 2 or 3.
One machine is fine.

Follow the steps on the page, in
order.

Take the screenshots the page asks
for.

**Answer the writeup items as you
go.**

Recap: what did you just see?

DISCUSS IN GROUPS

1

Mini-ImageNet

Did the ImageNet AI get anything wrong, right, and why?

2

Mini-AlphaGo

Nobody told it a strategy. Did it get better? How?

3

The AI agent, no human

When the robot got out of the sandbox, was it trying to escape? What was it actually doing?

4


The AI agent, regulation? Safety?

In your runs, what kept the robot in: a human saying no, a human on the kill switch, or a patched wall? If you were writing a rule for real AI labs, which one would you require, and who would check that they did it?

5

AI agents and you

Would you trust an AI agent to act for you, like booking travel or spending money? What would it need to have first?



BREAK • 5 MINUTES

Back in five

When you come back: the Gilman Street rules, then your side of the room is your side of the debate.

Debate rules: the Gilman Street rules

APPLY THESE TODAY

Same rules as Week 1, borrowed from 924 Gilman Street. Here's what each one means in a debate.

NO RACISM · NO SEXISM · NO HOMOPHOBIA · NO TRANSPHOBIA

Argue with ideas, never at people.

NO FIGHTING

No attacking a position without evidence. "That's wrong" isn't a point. "That's wrong, because Hugging Face found it first" is.

NO STAGEDIVING

No interrupting, piling on, or ambushing. One speaker at a time; let them finish. Respond to what they said, not to the person.

NO MAJOR LABEL BANDS

No blind love or hate for an AI company. OpenAI, Anthropic, Hugging Face: give reasons, not fandom.

The accord, layer by layer

SIGNED 9/29

Signed 9/29 by AI company leaders, including OpenAI and Anthropic. Voluntary. Trump called it “morally binding.”

LAYER 1

Internal controls

It says: Watch our models for dangerous skills (cyber, bio, chemical) and for reaching systems they shouldn't.

What it means: The company monitors its own AI. This is the layer that should catch something like the OpenAI and Hugging Face case.

Who checks: The company

LAYER 2

An internal team

It says: Give a team the job of making sure those controls work, and fixing what doesn't.

What it means: Someone inside checks the checkers. Still the same company.

Who checks: The company

LAYER 3

An outside auditor

It says: Partner with an independent auditor or evaluator to test whether the controls work.

What it means: The first outsider. The accord doesn't say who picks or pays them, how often, or whether results go public.

Who checks: An outsider the company partners with

LAYER 4

A board committee

It says: An independent board committee gets the reports and makes sure problems get fixed.

What it means: Problems go up to the company's own board, not out to the public or a regulator.

Who checks: The company's own board of directors

The debate

35 MIN

Topic: would the four-layer accord have caught the OpenAI and Hugging Face case?

SIDE A

It's enough, or a real start

Four layers, an outside auditor, and the FTC can hold companies to public promises. Faster than a law.

Cost to name: the companies wrote it, and the companies check it.

SIDE B

It doesn't go far enough

Reports stop at each company's own board. Nobody outside has to be told, and any penalty depends on the FTC choosing to act.

Cost to name: a law is slower, and a reporting clock can slow fixes.

How the debate runs: prep

16 MIN

1

Your side is where you sit

2 min

Left half of the room is Side A, right half is Side B. Get into a group of 2 or 3 on your side, then count off 1 to 4: that's your layer.

2

Research, with Claude

10 min

Your group gets one layer. Take it apart with Claude, following the steps on the next slide. Then write your notes in your own words.

3

Team up and pick speakers

4 min

Join one or two other groups on your side, so your team is 6 or 7 people: three teams per side. Pick your single best point, an opener, and a rebutter.

YOUR NOTES

1. Your strongest point, in one sentence
2. The evidence behind it
3. The honest cost of your position
4. What the other side will probably say

THE RULE ON AI

Use it to research. Don't read from it.

Step 2: use Claude to take your layer apart

1

Open Claude. Paste your layer's exact words from the full text (linked in the lab, section 5), since Claude may not know the accord exists. Say which side you're on.

2

Ask: "What does this mean in practice? What would a company actually have to do?"

3

Ask: "Who benefits from this, and what are its limits? What could still slip past it?"

4

Ask: "Who enforces this, and what happens if a company doesn't do it?" Then: "Has that ever actually happened?"

5

Ask: "Would this layer have caught the OpenAI and Hugging Face case? Who finds out, and when?"

Last: "What are the strongest points for my side, and the best objection from the other side?" Then close Claude and write your notes. Sample prompt: lab, section 5.

How the debate runs: the rounds

17 MIN

Round 1

Openings

5 min

Openers stand up where you are: A1, A2, A3, then B1, B2, B3. About 45 seconds each: your claim in one sentence, your evidence, and the cost of your plan.

Huddle

1 min

Back to your group. Rebutters pick one point the other side made.

Round 2

Rebuttals

5 min

Rebutters stand up and face the other side. Side B first this time, then A. Answer one specific point, with evidence. No new arguments.

Round 3

Switch

3 min

30 seconds with your team: what was the other side's best point? Then someone new argues it in one sentence.

Recap

What did you hear?

3 min

Which argument from the other side was hardest to answer? Then we're done.

THE GILMAN RULES, APPLIED

One speaker at a time. Let them finish.

Every point needs evidence.

Answer the argument, not the person.

Reasons, not fandom.

No winner is declared. The point is to hear the best case on both sides.

BEFORE YOU LEAVE

Upload your notes

Attendance = your notes from today, uploaded to Canvas right after class.

At home: lab sections 6 and 7 (discussion, and the five questions on your own AI).

Friday 10/2, 11:59 p.m.: Lab 1 writeup and the Chapter 1 quiz, both on Canvas.